

# N11736089\_VarunVikasJaiswal\_IFN695

June 1, 2025

```
[1]: # Load Required Libraries
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
from sklearn.ensemble import IsolationForest
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, LSTM, RepeatVector, TimeDistributed, Dense
import matplotlib.pyplot as plt

# Load the Dataset
df = pd.read_csv("C:/Users/Varun/OneDrive/Desktop/DATASET/merged_cleaned_data.csv", parse_dates=['SETTLEMENTDATE'])

# Feature Engineering
# Ensure datetime conversion
df['SETTLEMENTDATE'] = pd.to_datetime(df['SETTLEMENTDATE'], errors='coerce')

# Create lagged price
df['lagged_RRP'] = df['RRP'].shift(1).bfill()

# Calculate percentage price change
df['price_change_pct'] = df['RRP'].pct_change().fillna(0) * 100

# Calculate demand forecast error
df['demand_error'] = df['TOTALDEMAND'] - df['DEMANDFORECAST']

# Extract temporal features
df['hour_of_day'] = df['SETTLEMENTDATE'].dt.hour
df['day_of_week'] = df['SETTLEMENTDATE'].dt.dayofweek

# Define Selected Features
features = [
    'RRP', 'lagged_RRP', 'price_change_pct', 'demand_error',
    'TOTALDEMAND', 'AVAILABLEGENERATION', 'DISPATCHABLEGENERATION',
    'TRADING_RRP', 'TEMPERATURE', 'SOLAR_RADIATION', 'RAINFALL',
    'hour_of_day', 'day_of_week'
```

```
]

```

```
[2]: df

```

```
[2]:
      SETTLEMENTDATE REGIONID      RRP  MARKETSUSPENDEDFLAG \
0      2022-01-01 00:05:00    QLD1  118.73008              0.0
1      2022-01-01 00:10:00    QLD1  119.10000              0.0
2      2022-01-01 00:15:00    QLD1  118.73008              0.0
3      2022-01-01 00:20:00    QLD1  109.20000              0.0
4      2022-01-01 00:25:00    QLD1  117.61542              0.0
...
220660 2023-12-31 23:40:00    QLD1   85.55000              0.0
220661 2023-12-31 23:45:00    QLD1   85.55000              0.0
220662 2023-12-31 23:50:00    QLD1   85.75000              0.0
220663 2023-12-31 23:55:00    QLD1   85.55000              0.0
220664 2024-01-01 00:00:00    QLD1   85.55000              0.0

      RAISE6SECRP  RAISE60SECRP  RAISE5MINRRP  RAISEREGRRP  LOWER6SECRP \
0              5.00          2.29          0.75          19.94          0.18
1              5.00          2.29          0.75          19.94          0.18
2              5.00          1.98          0.75          19.94          0.18
3              5.00          2.29          0.75          19.94          0.18
4              5.00          2.29          0.75          19.94          0.18
...
220660          0.39          0.39          0.39          1.00          0.01
220661          0.39          0.39          0.39          1.00          0.03
220662          0.39          0.39          0.39          2.64          0.01
220663          0.39          0.39          0.38          0.91          0.03
220664          0.39          0.39          0.39          1.17          0.03

      LOWER60SECRP  ...  TRADING_RRP      DATE  RAINFALL  TEMPERATURE \
0              0.90  ...      118.73  2022-01-01        NaN          26.6
1              0.90  ...      119.10  2022-01-01        NaN          26.6
2              0.90  ...      118.73  2022-01-01        NaN          26.6
3              0.90  ...      109.20  2022-01-01        NaN          26.6
4              0.90  ...      117.62  2022-01-01        NaN          26.6
...
220660          0.07  ...       85.55  2023-12-31         5.8          30.4
220661          0.17  ...       85.55  2023-12-31         5.8          30.4
220662          0.08  ...       85.75  2023-12-31         5.8          30.4
220663          0.15  ...       85.55  2023-12-31         5.8          30.4
220664          0.17  ...       85.55  2024-01-01        NaN          NaN

      SOLAR_RADIATION  lagged_RRP  price_change_pct  demand_error \
0              10.1    118.73008          0.000000    5994.84561
1              10.1    118.73008          0.311564    6040.45600
2              10.1    119.10000         -0.310596    6025.94813

```

3	10.1	118.73008	-8.026677	5942.65281
4	10.1	109.20000	7.706429	5924.75816
...	...	...	...	...
220660	28.3	85.55000	0.000000	6438.64000
220661	28.3	85.55000	0.000000	6391.79000
220662	28.3	85.55000	0.233781	6399.35000
220663	28.3	85.75000	-0.233236	6343.06000
220664	NaN	85.55000	0.000000	6344.38000

	hour_of_day	day_of_week
0	0	5
1	0	5
2	0	5
3	0	5
4	0	5
...	...	...
220660	23	6
220661	23	6
220662	23	6
220663	23	6
220664	0	0

[220665 rows x 27 columns]

```
[3]: df.columns
```

```
[3]: Index(['SETTLEMENTDATE', 'REGIONID', 'RRP', 'MARKETSUSPENDEDFLAG',
        'RAISE6SECRRP', 'RAISE60SECRRP', 'RAISE5MINRRP', 'RAISEREGRRP',
        'LOWER6SECRRP', 'LOWER60SECRRP', 'LOWER5MINRRP', 'LOWERREGRRP',
        'TOTALDEMAND', 'DEMANDFORECAST', 'AVAILABLEGENERATION',
        'DISPATCHABLEGENERATION', 'NETINTERCHANGE', 'TRADING_RRP', 'DATE',
        'RAINFALL', 'TEMPERATURE', 'SOLAR_RADIATION', 'lagged_RRP',
        'price_change_pct', 'demand_error', 'hour_of_day', 'day_of_week'],
        dtype='object')
```

```
[4]: # Drop NaNs, handle inf values, and scale the features
X = df[features].replace([np.inf, -np.inf], np.nan).dropna()
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)
```

```
[8]: # -----
# Isolation Forest Anomaly Detection
# -----
# Train Isolation Forest on full scaled feature set
iso_model = IsolationForest(n_estimators=100, contamination=0.02,
    random_state=42)
df_iso = X.copy()
```

```
df_iso['iso_anomaly'] = iso_model.fit_predict(X_scaled)
# Convert output: -1 = anomaly + 1, 1 = normal + 0
df_iso['iso_anomaly'] = df_iso['iso_anomaly'].map({1: 0, -1: 1})
```

```
[9]: df_iso
```

```
[9]:
```

	RRP	lagged_RRP	price_change_pct	demand_error	TOTALDEMAND	\
575	91.20000	91.20000	0.000000	5963.68545	5931.35	
576	109.20000	91.20000	19.736842	5963.01764	5939.19	
577	109.09998	109.20000	-0.091593	5902.69941	5875.05	
578	107.42242	109.09998	-1.537635	5872.74354	5851.37	
579	109.10000	107.42242	1.561667	5888.37908	5853.82	
...	...	...	...	...	...	
220659	85.55000	60.60896	41.150747	6466.49000	6441.49	
220660	85.55000	85.55000	0.000000	6438.64000	6384.64	
220661	85.55000	85.55000	0.000000	6391.79000	6374.79	
220662	85.75000	85.55000	0.233781	6399.35000	6386.35	
220663	85.55000	85.75000	-0.233236	6343.06000	6296.06	
	AVAILABLEGENERATION	DISPATCHABLEGENERATION	TRADING_RRP	TEMPERATURE	\	
575	9525.71501	5733.61	91.20	30.2		
576	9529.07401	5622.78	109.20	30.2		
577	9523.92301	5528.22	109.10	30.2		
578	9524.36201	5500.82	107.42	30.2		
579	9533.37801	5494.25	109.10	30.2		
...	...	...	...	...		
220659	9324.73916	5912.86	85.55	30.4		
220660	9315.29446	5857.84	85.55	30.4		
220661	9115.39488	5848.00	85.55	30.4		
220662	9126.19154	5858.87	85.75	30.4		
220663	9388.81190	5767.10	85.55	30.4		
	SOLAR_RADIATION	RAINFALL	hour_of_day	day_of_week	iso_anomaly	
575	23.4	0.0	0	0	0	
576	23.4	0.0	0	0	0	
577	23.4	0.0	0	0	0	
578	23.4	0.0	0	0	0	
579	23.4	0.0	0	0	0	
...	...	...	...	...	...	
220659	28.3	5.8	23	6	0	
220660	28.3	5.8	23	6	0	
220661	28.3	5.8	23	6	0	
220662	28.3	5.8	23	6	0	
220663	28.3	5.8	23	6	0	

```
[217122 rows x 14 columns]
```

```
[10]: df_iso.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 217122 entries, 575 to 220663
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   RRP                                    217122 non-null float64
1   lagged_RRP                            217122 non-null float64
2   price_change_pct                      217122 non-null float64
3   demand_error                         217122 non-null float64
4   TOTALDEMAND                          217122 non-null float64
5   AVAILABLEGENERATION                  217122 non-null float64
6   DISPATCHABLEGENERATION               217122 non-null float64
7   TRADING_RRP                          217122 non-null float64
8   TEMPERATURE                          217122 non-null float64
9   SOLAR_RADIATION                      217122 non-null float64
10  RAINFALL                             217122 non-null float64
11  hour_of_day                          217122 non-null int32
12  day_of_week                          217122 non-null int32
13  iso_anomaly                          217122 non-null int64
dtypes: float64(11), int32(2), int64(1)
memory usage: 23.2 MB
```

```
[11]: # Count and print number of anomalies detected by Isolation Forest
num_iso_anomalies = df_iso['iso_anomaly'].sum()/len(df_iso)
print(f"Number of anomalies detected by Isolation Forest: {num_iso_anomalies}")
```

Number of anomalies detected by Isolation Forest: 0.0200025791951069

```
[12]: # Create Sliding Windows for LSTM
window_size = 96
X_seq = []
for i in range(len(X_scaled) - window_size):
    X_seq.append(X_scaled[i:i+window_size])
X_seq = np.array(X_seq)

# Build LSTM Autoencoder Model
input_dim = X_seq.shape[2]
inputs = Input(shape=(window_size, input_dim))
encoded = LSTM(64, return_sequences=False)(inputs)
bottleneck = RepeatVector(window_size)(encoded)
decoded = LSTM(64, return_sequences=True)(bottleneck)
outputs = TimeDistributed(Dense(input_dim))(decoded)

model = Model(inputs, outputs)
model.compile(optimizer='adam', loss='mse')
```

```
# Train the LSTM Autoencoder
```

```
model.fit(X_seq, X_seq, epochs=30, batch_size=64, validation_split=0.1,  
↳ verbose=1)
```

Epoch 1/30

3052/3052 329s 106ms/step

- loss: 0.0088 - val\_loss: 0.0013

Epoch 2/30

3052/3052 301s 99ms/step

- loss: 0.0010 - val\_loss: 8.6527e-04

Epoch 3/30

3052/3052 249s 82ms/step

- loss: 8.4124e-04 - val\_loss: 8.1116e-04

Epoch 4/30

3052/3052 265s 87ms/step

- loss: 6.1090e-04 - val\_loss: 5.9193e-04

Epoch 5/30

3052/3052 267s 87ms/step

- loss: 5.2574e-04 - val\_loss: 5.4852e-04

Epoch 6/30

3052/3052 271s 89ms/step

- loss: 4.6558e-04 - val\_loss: 4.5584e-04

Epoch 7/30

3052/3052 255s 84ms/step

- loss: 4.3536e-04 - val\_loss: 3.9063e-04

Epoch 8/30

3052/3052 257s 84ms/step

- loss: 4.0428e-04 - val\_loss: 3.8401e-04

Epoch 9/30

3052/3052 253s 83ms/step

- loss: 3.6473e-04 - val\_loss: 3.7030e-04

Epoch 10/30

3052/3052 263s 86ms/step

- loss: 3.5127e-04 - val\_loss: 4.2261e-04

Epoch 11/30

3052/3052 263s 86ms/step

- loss: 3.3407e-04 - val\_loss: 3.4077e-04

Epoch 12/30

3052/3052 245s 80ms/step

- loss: 3.1965e-04 - val\_loss: 3.1861e-04

Epoch 13/30

3052/3052 246s 80ms/step

- loss: 3.1448e-04 - val\_loss: 3.4746e-04

Epoch 14/30

3052/3052 247s 81ms/step

- loss: 3.0350e-04 - val\_loss: 2.9992e-04

Epoch 15/30

3052/3052 246s 81ms/step

```

- loss: 2.9553e-04 - val_loss: 3.2844e-04
Epoch 16/30
3052/3052          247s 81ms/step
- loss: 3.0602e-04 - val_loss: 3.1831e-04
Epoch 17/30
3052/3052          246s 81ms/step
- loss: 3.1669e-04 - val_loss: 3.0301e-04
Epoch 18/30
3052/3052          247s 81ms/step
- loss: 2.8562e-04 - val_loss: 2.6665e-04
Epoch 19/30
3052/3052          246s 81ms/step
- loss: 2.6990e-04 - val_loss: 4.7265e-04
Epoch 20/30
3052/3052          267s 88ms/step
- loss: 2.7615e-04 - val_loss: 2.5319e-04
Epoch 21/30
3052/3052          272s 89ms/step
- loss: 2.7147e-04 - val_loss: 2.8716e-04
Epoch 22/30
3052/3052          250s 82ms/step
- loss: 2.7016e-04 - val_loss: 2.3558e-04
Epoch 23/30
3052/3052          250s 82ms/step
- loss: 3.1250e-04 - val_loss: 2.4728e-04
Epoch 24/30
3052/3052          248s 81ms/step
- loss: 2.7019e-04 - val_loss: 2.8995e-04
Epoch 25/30
3052/3052          269s 83ms/step
- loss: 2.5757e-04 - val_loss: 2.5405e-04
Epoch 26/30
3052/3052          252s 83ms/step
- loss: 2.4842e-04 - val_loss: 2.4629e-04
Epoch 27/30
3052/3052          251s 82ms/step
- loss: 2.3661e-04 - val_loss: 2.2292e-04
Epoch 28/30
3052/3052          251s 82ms/step
- loss: 3.1688e-04 - val_loss: 4.3981e-04
Epoch 29/30
3052/3052          251s 82ms/step
- loss: 3.0490e-04 - val_loss: 3.6510e-04
Epoch 30/30
3052/3052          252s 83ms/step
- loss: 2.6952e-04 - val_loss: 2.5410e-04

```

```
[12]: <keras.src.callbacks.history.History at 0x20b6f100d50>
```

```
[13]: # Reconstruct and Compute Reconstruction Error
X_pred = model.predict(X_seq)
reconstruction_errors = np.mean(np.square(X_seq - X_pred), axis=(1, 2))

# Set Threshold and Flag Anomalies
threshold = np.percentile(reconstruction_errors, 95)
anomalies = reconstruction_errors > threshold

# Attach LSTM Anomaly Labels to Data
anomaly_df = X.iloc>window_size:].copy()
anomaly_df['reconstruction_error'] = reconstruction_errors
anomaly_df['lstm_anomaly'] = anomalies.astype(int)

# Combine Both Anomaly Flags
anomaly_df['iso_anomaly'] = df_iso.iloc>window_size:].['iso_anomaly'].values
```

6783/6783                      124s 18ms/step

```
[14]: anomaly_df.tail(10)
```

```
[14]:
```

	RRP	lagged_RRP	price_change_pct	demand_error	TOTALDEMAND	\
220654	65.47628	61.87869	5.813940	6616.48	6589.48	
220655	64.98145	65.47628	-0.755739	6574.99	6575.99	
220656	65.45639	64.98145	0.730886	6613.07	6601.07	
220657	65.35098	65.45639	-0.161039	6582.44	6555.44	
220658	60.60896	65.35098	-7.256234	6510.98	6492.98	
220659	85.55000	60.60896	41.150747	6466.49	6441.49	
220660	85.55000	85.55000	0.000000	6438.64	6384.64	
220661	85.55000	85.55000	0.000000	6391.79	6374.79	
220662	85.75000	85.55000	0.233781	6399.35	6386.35	
220663	85.55000	85.75000	-0.233236	6343.06	6296.06	

	AVAILABLEGENERATION	DISPATCHABLEGENERATION	TRADING_RRP	TEMPERATURE	\
220654	9570.58009	6127.72	65.48	30.4	
220655	9579.74810	6135.83	64.98	30.4	
220656	9575.39208	6132.63	65.46	30.4	
220657	9568.80508	6125.14	65.35	30.4	
220658	9581.54185	6187.67	60.61	30.4	
220659	9324.73916	5912.86	85.55	30.4	
220660	9315.29446	5857.84	85.55	30.4	
220661	9115.39488	5848.00	85.55	30.4	
220662	9126.19154	5858.87	85.75	30.4	
220663	9388.81190	5767.10	85.55	30.4	

	SOLAR_RADIATION	RAINFALL	hour_of_day	day_of_week	\
220654	28.3	5.8	23	6	



220655	28.3	5.8	23	6
220656	28.3	5.8	23	6
220657	28.3	5.8	23	6
220658	28.3	5.8	23	6
220659	28.3	5.8	23	6
220660	28.3	5.8	23	6
220661	28.3	5.8	23	6
220662	28.3	5.8	23	6
220663	28.3	5.8	23	6

	reconstruction_error	lstm_anomaly	iso_anomaly
220654	0.000115	0	0
220655	0.000110	0	0
220656	0.000108	0	0
220657	0.000107	0	0
220658	0.000105	0	0
220659	0.000105	0	0
220660	0.000105	0	0
220661	0.000106	0	0
220662	0.000110	0	0
220663	0.000114	0	0

```
[15]: # Print anomaly percentages from LSTM and Isolation Forest
lstm_count = anomaly_df['lstm_anomaly'].sum()
lstm_percent = lstm_count / len(anomaly_df) * 100

iso_count = anomaly_df['iso_anomaly'].sum()
iso_percent = iso_count / len(anomaly_df) * 100

print(f"LSTM Anomalies: {lstm_count} ({lstm_percent:.2f}%)")
print(f"Isolation Forest Anomalies: {iso_count} ({iso_percent:.2f}%)")
```

LSTM Anomalies: 10852 (5.00%)  
Isolation Forest Anomalies: 4343 (2.00%)


```
[16]: # Compare agreement between LSTM and Isolation Forest
anomaly_df['both_agree'] = (anomaly_df['lstm_anomaly'] == 1) &
    (anomaly_df['iso_anomaly'] == 1)
both_agree_count = anomaly_df['both_agree'].sum()
print(f"Both models agree on {both_agree_count} anomalies ({(both_agree_count/
    len(anomaly_df))*100:.2f}%)")
```

Both models agree on 1563 anomalies (0.72%)

```
[17]: # Save final combined anomaly output for future use
anomaly_df.to_csv("anomaly_detected_combined.csv", index=False)
```

```
[18]: from matplotlib_venn import venn2
from matplotlib import pyplot as plt

plt.figure(figsize=(6,6))
venn2(subsets = (
    ((anomaly_df['lstm_anomaly'] == 1) & (anomaly_df['iso_anomaly'] == 0)).
    ↪sum(),
    ((anomaly_df['iso_anomaly'] == 1) & (anomaly_df['lstm_anomaly'] == 0)).
    ↪sum(),
    ((anomaly_df['lstm_anomaly'] == 1) & (anomaly_df['iso_anomaly'] == 1)).sum()
), set_labels=('LSTM', 'Isolation Forest'))
plt.title("Overlap of Detected Anomalies")
plt.show()
```



```
[19]: # Top anomalies by LSTM reconstruction error
top_anomalies = anomaly_df.sort_values(by='reconstruction_error',
    ↪ascending=False).head(10)
print(top_anomalies[['RRP', 'reconstruction_error', 'lstm_anomaly',
    ↪'iso_anomaly']])
```

	RRP	reconstruction_error	lstm_anomaly	iso_anomaly
9237	70.81522	0.030475	1	0
9236	70.81522	0.030418	1	1

9238	-999.99417	0.030147	1	0
9235	-999.99171	0.030082	1	1
9239	-999.99417	0.029818	1	1
9234	-999.99171	0.029789	1	0
9233	86.04854	0.029560	1	0
9240	300.10000	0.029533	1	1
9232	86.04854	0.029168	1	0
9241	300.10000	0.029159	1	0

```
[20]: # Check for RRP values equal to -999.99 in the original dataset
invalid_rrp_rows = df[df['RRP'] == -999.99]

# Display results
print(f"Number of rows with RRP = -999.99: {len(invalid_rrp_rows)}")
invalid_rrp_rows.head()
```

Number of rows with RRP = -999.99: 0

```
[20]: Empty DataFrame
Columns: [SETTLEMENTDATE, REGIONID, RRP, MARKETSUSPENDEDFLAG, RAISE6SECRRP,
RAISE60SECRRP, RAISE5MINRRP, RAISEREGRRP, LOWER6SECRRP, LOWER60SECRRP,
LOWER5MINRRP, LOWERREGRRP, TOTALDEMAND, DEMANDFORECAST, AVAILABLEGENERATION,
DISPATCHABLEGENERATION, NETINTERCHANGE, TRADING_RRP, DATE, RAINFALL,
TEMPERATURE, SOLAR_RADIATION, month]
Index: []

[0 rows x 23 columns]
```

```
[21]: # Load the dataset
df = pd.read_csv("anomaly_detected_combined.csv")

# Create 'both_agree' column if not present
df['both_agree'] = (df['lstm_anomaly'] == 1) & (df['iso_anomaly'] == 1)

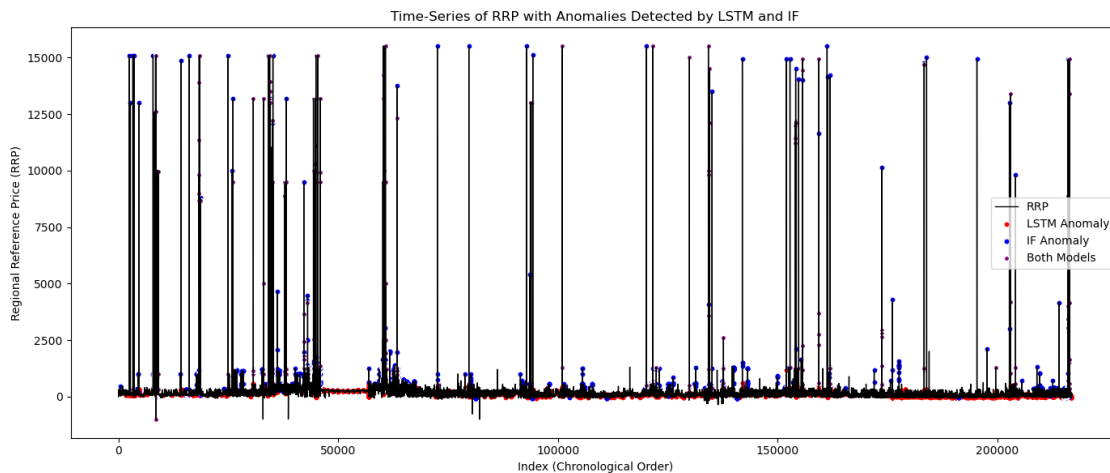
# Create index for plotting (since no timestamp column)
df['index'] = range(len(df))

# Plot RRP time-series
plt.figure(figsize=(14, 6))
plt.plot(df['index'], df['RRP'], label='RRP', color='black', linewidth=1)

# Overlay anomalies
plt.scatter(df[df['lstm_anomaly'] == 1]['index'], df[df['lstm_anomaly'] == 1]['RRP'],
            color='red', label='LSTM Anomaly', s=10)
plt.scatter(df[df['iso_anomaly'] == 1]['index'], df[df['iso_anomaly'] == 1]['RRP'],
            color='blue', label='IF Anomaly', s=10)
```

```
plt.scatter(df[df['both_agree'] == True]['index'], df[df['both_agree'] == True]['RRP'],
            color='purple', label='Both Models', s=20, edgecolor='white')

# Formatting
plt.title("Time-Series of RRP with Anomalies Detected by LSTM and IF")
plt.xlabel("Index (Chronological Order)")
plt.ylabel("Regional Reference Price (RRP)")
plt.legend()
plt.tight_layout()
plt.show()
```



```
[22]: # Load your data
df = pd.read_csv("anomaly_detected_combined.csv")

# Confirm both_agree column exists
df['both_agree'] = (df['lstm_anomaly'] == 1) & (df['iso_anomaly'] == 1)

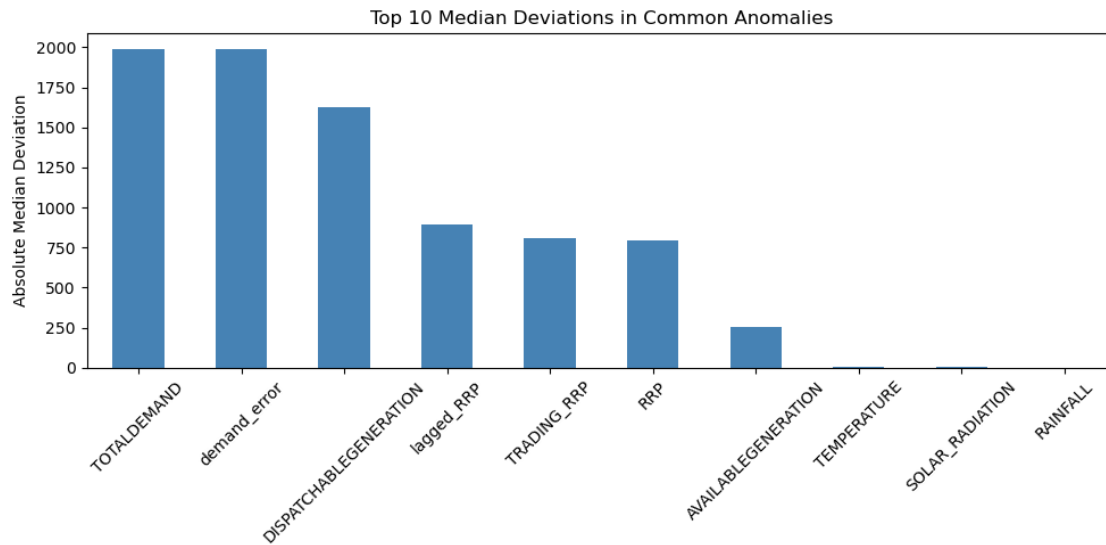
# Keep only original (unscaled) meaningful features
core_features = [
    'RRP', 'lagged_RRP', 'price_change_pct', 'demand_error', 'TOTALDEMAND',
    'AVAILABLEGENERATION', 'DISPATCHABLEGENERATION', 'TRADING_RRP',
    'TEMPERATURE', 'SOLAR_RADIATION', 'RAINFALL'
]

# Calculate median deviation
normal_median = df[df['both_agree'] == False][core_features].median()
anomaly_median = df[df['both_agree'] == True][core_features].median()
deviation = (anomaly_median - normal_median).abs().sort_values(ascending=False)
```

```

# Plot top 10
plt.figure(figsize=(10, 5))
deviation.head(10).plot(kind='bar', color='steelblue')
plt.title("Top 10 Median Deviations in Common Anomalies")
plt.ylabel("Absolute Median Deviation")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```

[23]: import pandas as pd
import matplotlib.pyplot as plt

# Load your anomaly detection results (no date column to parse)
df = pd.read_csv("anomaly_detected_combined.csv")

# Create the both_agree column
df['both_agree'] = (df['lstm_anomaly'] == 1) & (df['iso_anomaly'] == 1)

# Filter where both models agree
common_anomalies = df[df['both_agree'] == True]

# Sort by reconstruction error and pick top 10
top_common = common_anomalies.sort_values(by='reconstruction_error',
↪ascending=False).head(10)

# Display the top anomalies
print("Top 10 High-Confidence Anomalies (LSTM + IF):")

```

```

print(top_common[['RRP', 'reconstruction_error', 'TOTALDEMAND', 'TEMPERATURE',
↪ 'demand_error']])

# Visualize as a bar chart
plt.figure(figsize=(10, 5))
plt.barh(top_common.index.astype(str), top_common['reconstruction_error'],
↪ color='red')
plt.xlabel("Reconstruction Error")
plt.title("Top 10 Common Anomalies Detected by Both Models")
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

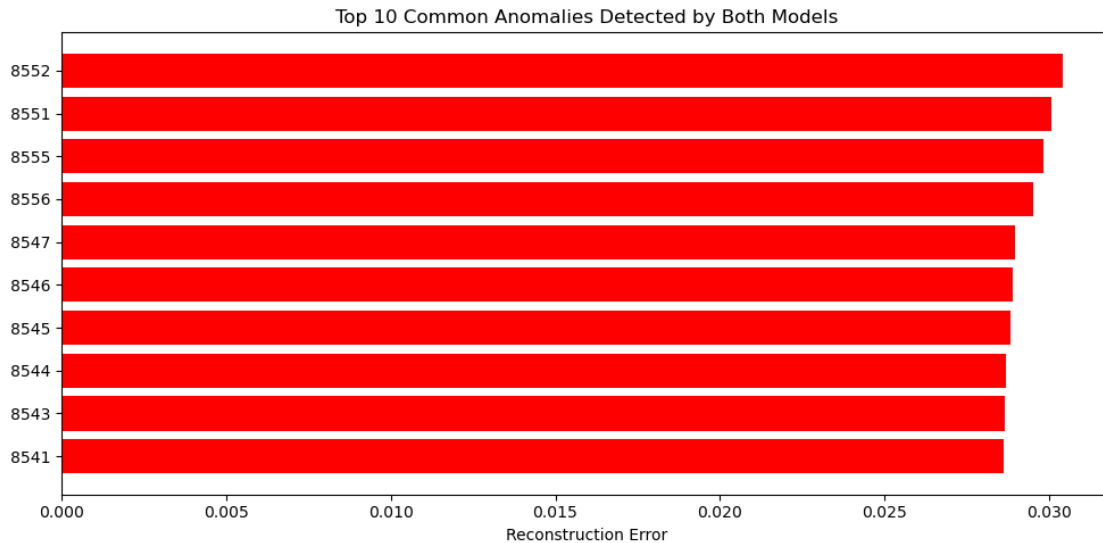
```

Top 10 High-Confidence Anomalies (LSTM + IF):

	RRP	reconstruction_error	TOTALDEMAND	TEMPERATURE \
8552	70.81522	0.030418	9068.35	34.7
8551	-999.99171	0.030082	9149.24	34.7
8555	-999.99417	0.029818	9068.35	34.7
8556	300.10000	0.029533	9086.19	34.7
8547	77.07020	0.028984	9268.48	34.7
8546	77.07020	0.028890	9268.48	34.7
8545	15100.00000	0.028829	9268.48	34.7
8544	15100.00000	0.028698	9268.48	34.7
8543	11.00008	0.028667	9205.84	34.7
8541	12594.00000	0.028637	9205.84	34.7

	demand_error
8552	9112.66738
8551	9157.56813
8555	9112.66738
8556	9119.83648
8547	9244.11867
8546	9244.11867
8545	9244.11867
8544	9244.11867
8543	9223.27262
8541	9223.27262




```
[24]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load anomaly file
df = pd.read_csv("anomaly_detected_combined.csv")

# Create time-based features if not already present
df['hour_of_day'] = df.get('hour_of_day', pd.Series(range(len(df))) % 24)
df['day_of_week'] = df.get('day_of_week', pd.Series(range(len(df))) % 7)

# Create a pivot table showing count of LSTM anomalies by hour and day
heatmap_data = df[df['lstm_anomaly'] == 1].pivot_table(
    index='day_of_week',
    columns='hour_of_day',
    values='RRP',
    aggfunc='count',
    fill_value=0
)

# Plot the heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(heatmap_data, cmap="Reds", annot=True, fmt='d')
plt.title("Heatmap of LSTM Anomalies by Day of Week and Hour of Day")
plt.xlabel("Hour of Day")
plt.ylabel("Day of Week (0 = Monday)")
plt.tight_layout()
plt.show()
```



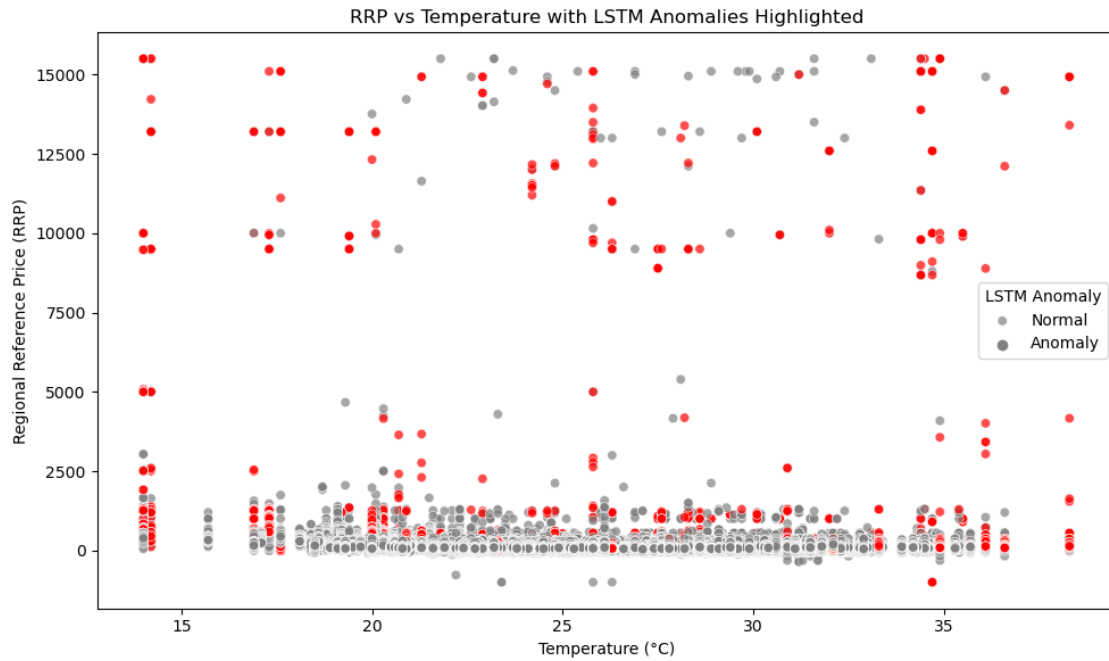
```
[25]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load your data
df = pd.read_csv("anomaly_detected_combined.csv")

# Scatter plot: RRP vs Temperature, colored by LSTM anomaly
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='TEMPERATURE', y='RRP', hue='lstm_anomaly',
                palette={0: 'gray', 1: 'red'}, alpha=0.7)

plt.title("RRP vs Temperature with LSTM Anomalies Highlighted")
plt.xlabel("Temperature (°C)")
plt.ylabel("Regional Reference Price (RRP)")
plt.legend(title="LSTM Anomaly", labels=["Normal", "Anomaly"])
plt.tight_layout()
plt.show()
```





[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	